



A gentle Introduction to deep learning with JAX

BAMB! Summer School
Tutorial 5


We will focus on **regression models**

From inputs to outputs

$$y = f(x, \theta)$$

Diagram illustrating the components of the regression equation $y = f(x, \theta)$:


- y : output
- f : function
- x : input
- θ : parameters




Learning by minimizing a loss

$$l(\theta) = \sum_n \|y_n - f(x, \theta)\|^2$$


Gradient Based Optimization




Gradient Based Optimization




Gradient Based Optimization



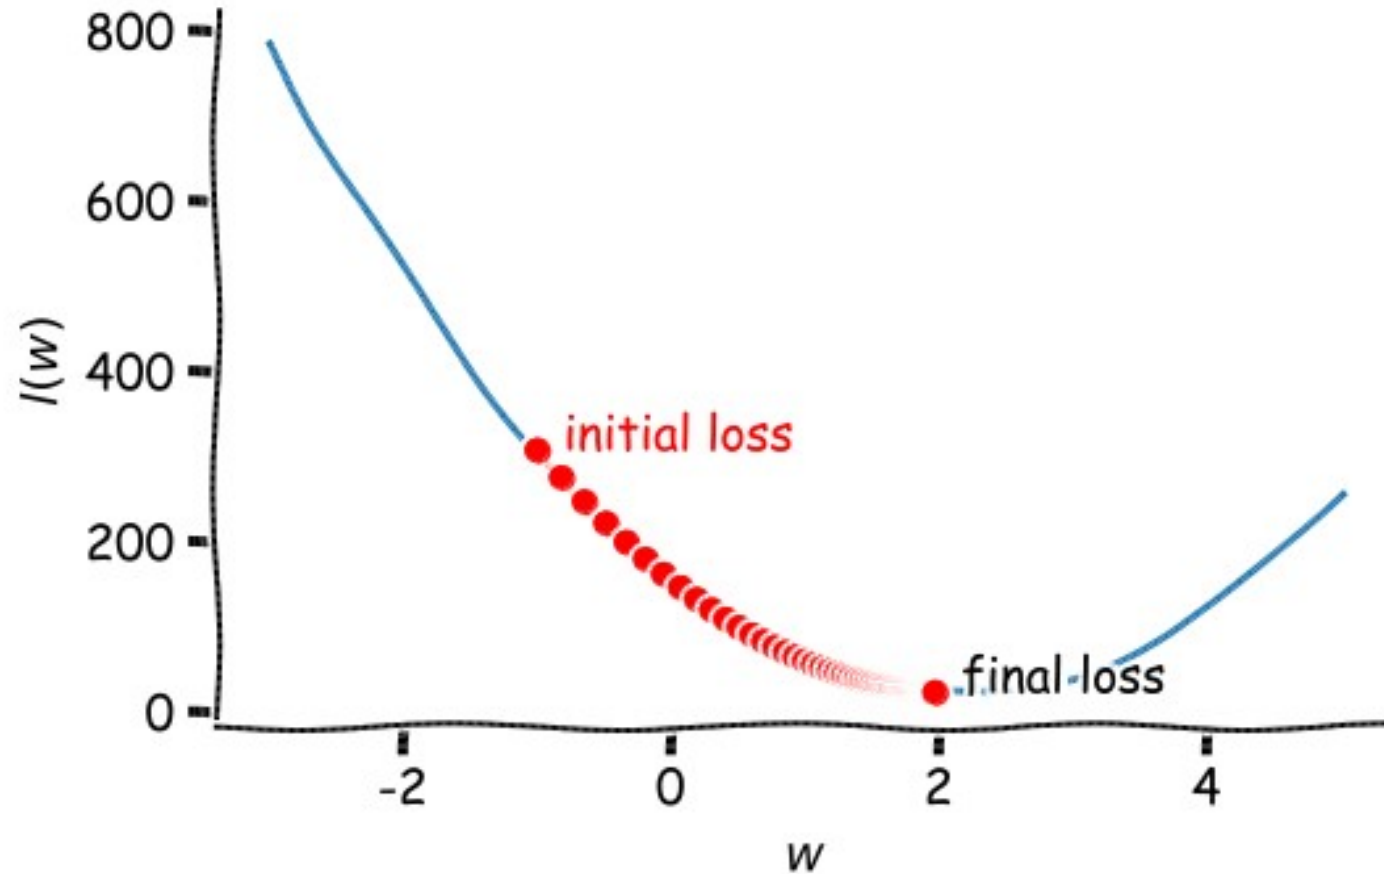
Gradient Based Optimization



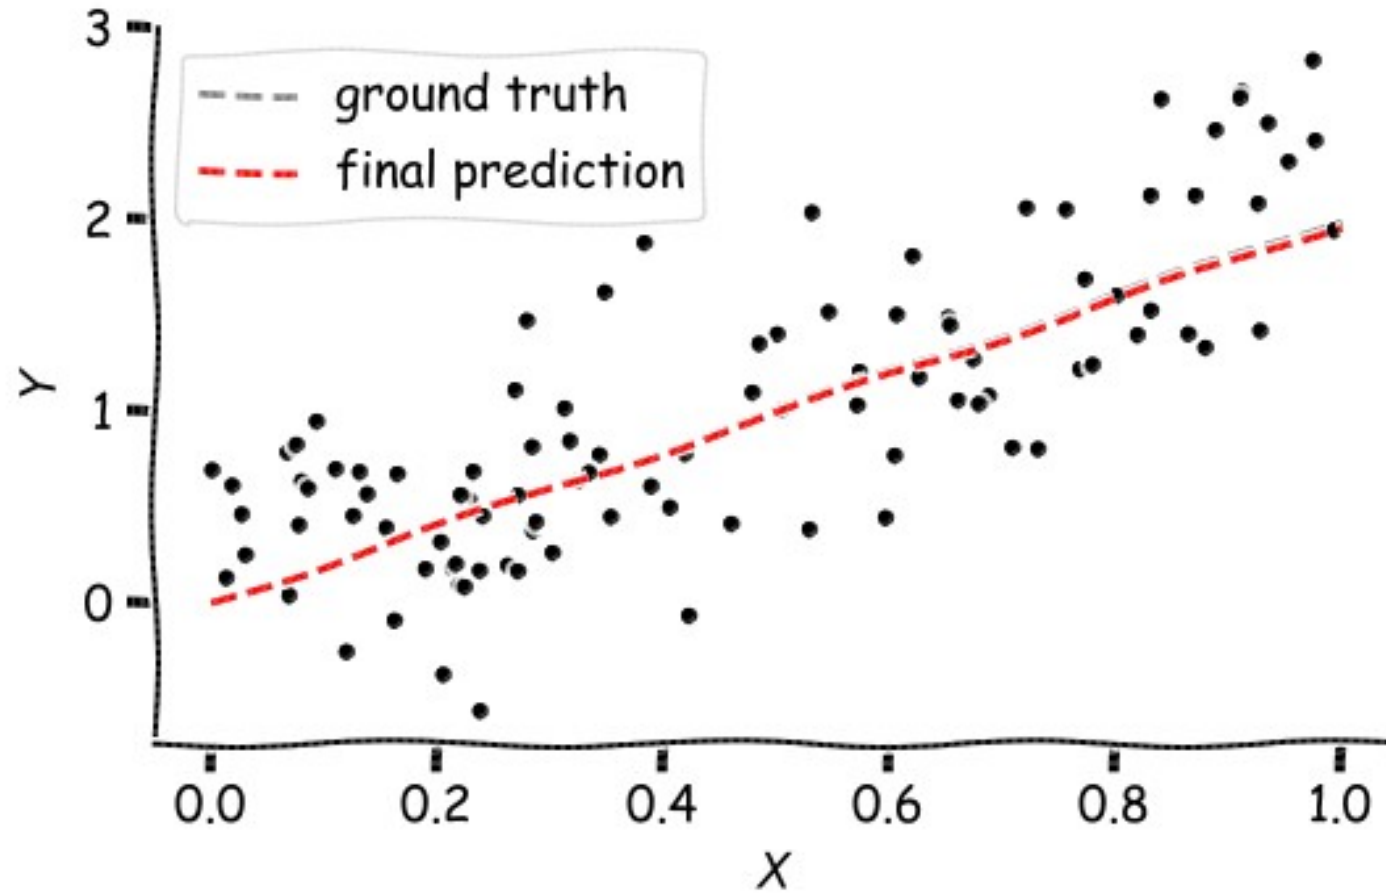
Gradient Based Optimization



Gradient Based Optimization




Gradient Based Optimization



We will work with different models


- Linear / Logistic Regression (as a warm up)
- Neural networks (NN) for image classification
- Convolutional neural networks (CNN)



Let's get started coding!

What is JAX for?

- Perform **Gradient Based** optimization ...
- ... in LARGE models (millions of parameters)
-



A bit about JAX

JAX is a library

- * in Python
- * ... made to be very close to the python library **numpy**
- * that computes gradient automatically for you (**automatic differentiation**)
- * that can compile code on demand to be very fast (**just in time** [or jit] compilation)
- * that has useful features to 'vectorize' function : **vmap**
(apply a function designed for scalars to vectors or tensors)

Some resources for the less experienced coders

A course on using python for scientific computing

<http://ucl-cs-grad.github.io/scipython/>

Numpy for matlab (by Numpy creators)

<https://numpy.org/doc/stable/user/numpy-for-matlab-users.html>

Numpy for matlab users cheat sheet:

https://mas-dse.github.io/DSE200/cheat_sheets/1_python/6_2_NumPy_for_MATLAB_users.pdf

Tutorial on github

https://github.com/vincentadam87/intro_to_jax